

<https://github.com/binpass/popl26-tutorial>

Clone this repo!

Working through src/solution.py

Easiest to work in docker

Safest too 😅

NB we're shooting for sound-ish

We'll use bash throughout, necessary for part 2



Analyzing Shell Scripts

POPL 2026 // Rennes, France

Tuesday, January 13th, 2026

Michael Greenberg, Konstantinos Kallas,
Nikos Vasilakis, and Evangelos Lamprou



Who we are



POPL 2020 Executable Formal Semantics for the POSIX Shell

EuroSys 2021 PaSh: light-touch data-parallel shell processing

HotOS 2021 Unix shell programming: the next 50 years

ICFP 2021 An order-aware dataflow model for parallel Unix pipelines

PPoPP 2022 Automatic synthesis of parallel unix commands and pipelines with KumQuat

OSDI 2022 Practically Correct, Just-in-Time Shell Script Parallelization

NSDI 2023 DiSh: Dynamic Shell-Script Distribution

HotOS 2023 Executing Shell Scripts in the Wrong Order, Correctly

ATC 2025 The Koala Benchmarks for the Shell: Characterization and Implications

HotOS 2025 From Ahead-of- to Just-in-Time and Back Again: Static Analysis for Unix Shell Programs



Who we are



BROWN



UCLA

Some systems papers:

- [NSDI 26, ATC 25 (best paper), NSDI 23, OSDI 22, EuroSys 21 (best paper), HotOS 25, HotOS 23, HotOS 21]

And some PL papers:

- [POPL 20] Executable Formal Semantics for the POSIX Shell
- [ICFP 21] An order-aware dataflow model for parallel Unix pipelines



Recent PaSh Retreat, Fall 25





Our work today



shell overview

what is even happening here?



parsing POSIX shell

finger exercises in Python with [libdash](#) 1, 2, 3



syntactic analysis

identifying and counting shell features 4



break



semantic analysis

identifying pure script fragments 5



stubs for dynamic work

replace pure fragments with stubs 6



JIT instrumentation

replace dangerous commands with [try](#) 7, 8

<https://github.com/binpass/popl26-tutorial>

Clone this repo!

Working through src/solution.py

Easiest to work in docker

Safest too 😅

NB we're shooting for sound-ish

We'll use bash throughout, necessary for part 2





Step 1: parsing POSIX shell (step1_parse_script)

Parse a shell script into a Python object AST!

shasta

json_to_ast.
to_ast_node

dash

libdash

parser.parse

The `parse_shell_to_asts` function is a generator---properly, the shell must be parsed line-by-line. We can fudge it and call `list`.


Run your code to see the AST!

utils.py

parse_shell_to_asts

solution.py

step1_parse_script





 Step 2: Creating a function that walks the AST

Create an AST visitor! We'll have it work in pre-order, with both visiting and replacement. You have two tasks:

1. Call `walk_ast` with `visit=print`
2. Implement the `AST.CommandNode` case of `utils.walk_ast`

Run your code on a few examples to see how the visitor works.






Step 3: Pretty-print the script

Render the AST back as a shell script!

Interposing on shell scripts means taking a script as input, manipulating it, and then producing a modified script as output.

Complete the `ast_to_code` function `utils.py`; it takes a list of `AstNodes` and returns a string. It won't look exactly the same... play spot the difference with a few scripts!






Step 4: Syntactic static analysis of a shell script

Let's write a simple syntactic analysis: how many subshells will a script create?

You might be wondering... which statements create subshells?

There are four ways to create a subshell:

1. asynchronous commands (`... &`; makes 1 subshell)
2. pipes (`... | ... | ...`; makes n or $n-1$ [depends on shell!])
3. subshells (`(...); 1`)
4. command substitutions
(`$(...); 1`)





break time!






Step 5: Static analysis

Let's build a more serious analysis: which commands can affect the shell's binding? You'll implement a predicate `is_effect_free` using `walk_ast`. What can affect bindings in the shell? We'll use the following:

- function definitions (`fun() { ... }`)
- commands with a non-zero number of assignments (`VAR=val ...`)
- assigning parameter expansion (`${VAR=val}` or `${VAR:=val}`)
- arithmetic expansion
`($(...))`

Some builtin commands do, too...
but let's not worry about that.





“Analyzing” shell scripts



```
cmd=rm; $cmd -rf /
```

```
if gcc -o mystery mystery.c; ./mystery
```

```
then cmd=rm; else cmd=echo; fi;
```


```
$cmd -rf /
```

```
gcc -o gen_code gen_code.c && ./gen_code >ohno.sh
```

```
. ohno.sh
```

**runtime is the
right time**

JIT architecture



Steps 6-8: JIT intervention for safe `rm`

We will work our way up to building a small JIT that interposes on invocations of `rm`.

1. **Step 6:** We'll use `walk_ast` to stub out certain commands---the core function of the preprocessor.
2. **Step 7:** We'll define a `jit.sh` that is invoked for each stub---avoiding repeating code during preprocessing.
3. **Step 8:** We'll have our JIT capture the shell's environment and hand it off an expansion routine (`expand.py`), which lets us dynamically detect invocations of `rm` that don't statically appear (e.g., `cmd=rm; $cmd -rf /`).




Step 6: stubs for dynamic interposition


Let's implement a 'dry run' transformation. We'll alter the script to print out each non-effectful command rather than running it.

Walk the AST with a `replace=...` function that finds nodes that satisfy `is_effect_free`.

For each such fragment, pretty print the code into a 'stub' in a known location, and then generate new code that `cats` that stub out, instead.




JIT architecture



⚠ Step 7: Build a simple JIT interposition

Build a JIT script that debugs effect-free fragments *and* runs them, like a minimal `set -x`. You'll need to fill in a bit of `replace_with_debug_jit` so that it generates the JIT invocation of the correct stub:


```
JIT_INPUT=/PATH/TO/STUB . src/debug_jit.sh
```



⚠ Step 8: Interpose on `rm`, just in time

Identify invocations of `rm` just in time, and wrap them in a semi-isolated container using `try`.

Our interposition JIT will identify every simple command invocation (with or without effects!) and stub it to hand off to the JIT, which captures the shell environment and hands it to `expand.py`, which can use `sh_expand` to determine which command is being invoked even for, e.g., `cmd=rm; $cmd -rf /`.



⚠ Step 8: Interpose on `rm`, just in time

The work here is complicated enough that you should do it in two steps.

1. Simply prepend `try` to each `CommandNode`.
2. Implement the optimization.
 - a. Uncomment the `deepcopy` call.
 - b. Expand the command using `expand.expand_command`.
 - c. Extract the executable name using `string_of_expanded_arg`.
 - d. If the executable is not unsafe per `unsafe_commands`, then you need not prepend `try` and can return early.

Where do we go from here?

Today: built a proof-of-concept JIT for dynamic intervention.

What moves it beyond a proof-of-concept? [PaSh](#)'s JIT jumps through more hoops:

env save/restore perf optimizations control (`break`, `continue`, `return`, .)

What's can *you* do with it?

your analysis here? your intervention here? can your work use [Koala](#)?

What can *we* do with what you have?

extend Koala? string reasoning? filesystem reasoning? binary analysis?

thank you!

feedback form at

<https://forms.gle/q87M5FzXSUwMZnFf6>

please let us know what you thought!

